AI Spend Cycles — Admin Guide

For 50pick staff who run the platform. No technical knowledge required. Everything here lives on one screen: **Admin → AI usage**.

What a cycle is

50pick pays Anthropic for every AI call it makes — writing polls, resolving markets, answering the help chatbot. A **cycle** is simply a fixed amount of that spend, so it can be counted.

> **One cycle = $100 of Claude spend.**
>
> Top up $1,000 and that is **10 cycles**. Top up $10,000 and that is **100 cycles**. The size is a setting — you can change it — but $100 is what 50pick uses today.

A cycle is **an amount, not a length of time**. It does not end on a Monday or at month end. It ends when $100 has been spent — which might take two weeks or two months, depending on how busy the platform is. **How long each one lasted is recorded**, so you can see that pattern.

Cycles are counted

Cycle 7 is always cycle 7. The numbering never resets, so "cycles this year" is a figure you can compare year to year.

Cycles are divisible

Because they are countable, they can be divided by the markets resolved — which is how the platform answers "what does one poll cost us?"

Cycles are a checkpoint

When one ends, the AI stops and waits for you. That is deliberate: it is the moment to look at the spend before more is committed.

The main panel

Admin → AI usage → Spend cycles. The four figures across the top are the whole story: how many cycles have finished, how full the current one is, how many your Anthropic credit has paid for, and the projected rate per year.

FIGURE	WHAT IT MEANS
Cycles closed	How many complete $100 cycles the platform has spent since it started. An exact count — not an estimate.
This cycle	How much of the current $100 has gone. The bar underneath shows the same thing. At 100% it closes and the AI pauses.
Funded cycles	How many cycles your Anthropic credit limit will pay for. **If this is below 1, your credit limit is smaller than one cycle** — see **Before you press Start**.
Cycles per year	The projected annual rate. Withheld until there is enough history — if it says "not enough data", that is the platform refusing to guess.

Reconciliation — the line at the bottom of the panel.
It compares the cycle ledger against the record of individual API calls. It should say **"They agree exactly."** If it ever reports a difference, the figures on this page should not be used for pricing until someone has looked into it.

When a cycle ends — the one thing you must act on

When the current cycle reaches $100 it closes, and **AI poll posting and AI market resolving stop**. A red bar appears at the top of the page:

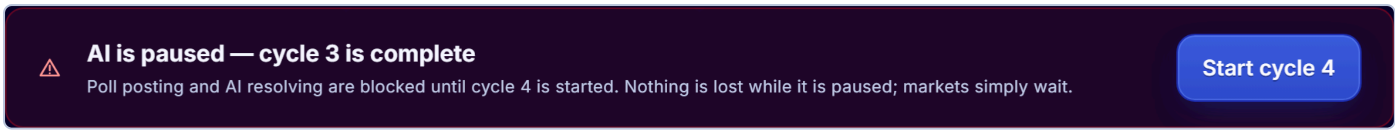

This is the bar to watch for. Nothing is broken and nothing is lost — markets simply wait. Press **Start cycle 4** to continue.

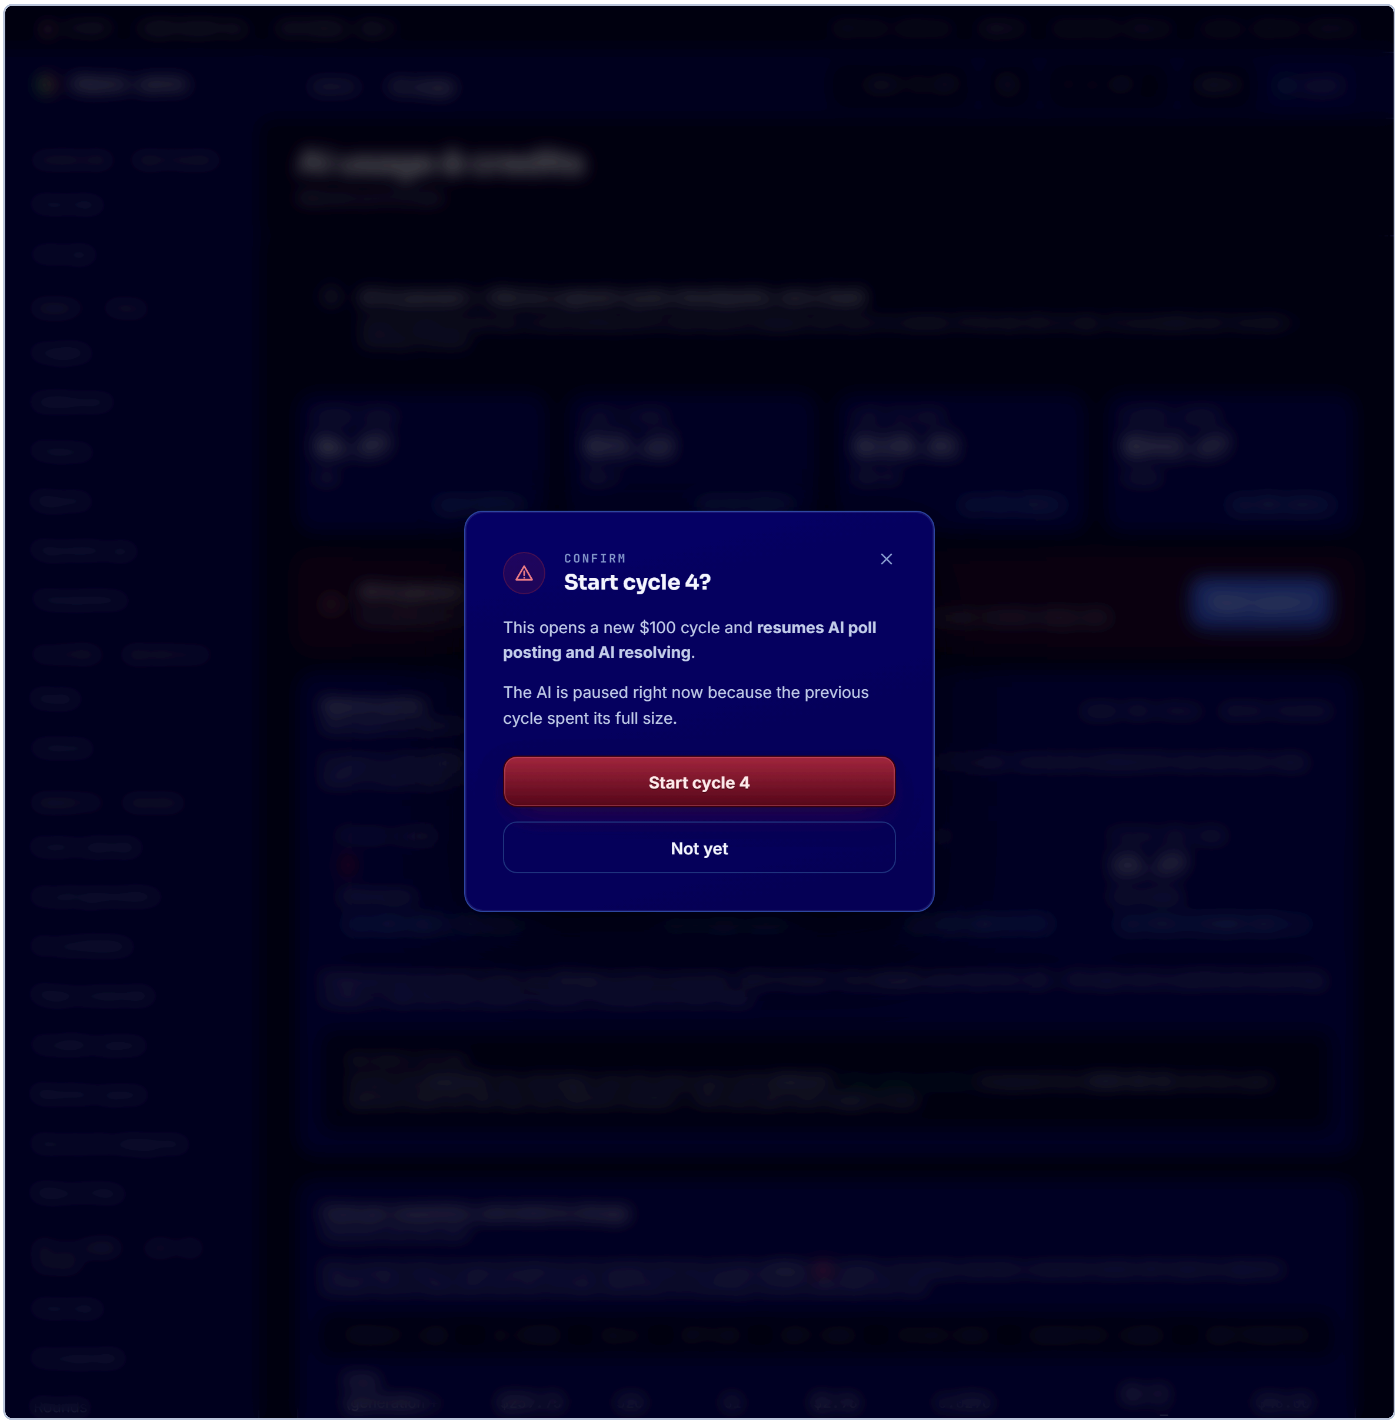

You are asked to confirm, because starting a cycle commits another $100 of spend. The platform records **who** started it.

Before you press Start, check two things.

1 · Is there still credit at Anthropic? A cycle is 50pick's own checkpoint. It does not know your Anthropic balance. If the balance is empty, starting a cycle will not help — top up first, then use **New top-up window** in the Credit budget panel below.

2 · Does the spend look normal? Compare how long this cycle lasted against the ones before it in **Every cycle**. A cycle that finished far quicker than usual is worth a question before you commit the next $100.

Closing a cycle early

The **Close cycle early** button ends the current cycle before it reaches $100 — useful for closing the books at the end of a month. **It also pauses the AI**, exactly as a natural ending does, and you then start the next one as normal.

⚠️ Closing early records a short cycle in the ledger. That is fine and it is honest — the yearly projection is worked out from **spend**, not from how many cycles are in the list, so closing early does not distort it.

The cycle history

Every cycle
Kila mzunguko

3 total

#	OPENED	CLOSED	LASTED	SIZE	SPENT	USED	RATES	STATUS
3	2026-08-14 19:57:39Z	–	–	$100.00	$42.67	43%	p3193cx	OPEN backfilled from usage history 2026-08-23
2	2026-07-20 20:26:27Z	2026-08-14 19:57:39Z	24d 23h	$100.00	$100.00	100%	p3193cx	CLOSED backfilled from usage history 2026-08-23
1	2026-06-26 19:57:39Z	2026-07-20 20:26:27Z	24d 0h	$100.00	$100.00	100%	p3193cx	CLOSED backfilled from usage history 2026-08-23

⛔ This ledger is **never pruned**. The per-call ledger below is kept 180 days; a cycle carries its own total so the count stays true long after the calls behind it are gone. "Used" above 100% means the final call crossed the boundary — real, bounded by one call, and shown rather than smoothed away.

Every cycle, oldest at the bottom. Lasted is the figure to watch — it tells you how quickly the platform is spending. **Size** is the amount that cycle was opened with, so changing the setting later never rewrites history.

In the example above, cycle 1 lasted 24 days and cycle 2 lasted 25 days. That is the platform's real burn rate: roughly **$100 every three and a half weeks**, so a $1,000 top-up would last about eight months.

Cycles by year
Mizunguko kwa mwaka

YEAR	CYCLES CLOSED	SPEND	AVG CYCLE LASTED	STATUS
2026	2	$200.00	24.5d	PART YEAR

Counted from the ledger, in the platform timezone — each closed cycle at the size it was **opened** with, never re-derived from today's size. A part year is marked, because a part-year total read as a full one understates the rate.

Cycles by year. The yearly total. A year still running is marked **PART YEAR**, so a partial total is never mistaken for a full one.

What a market costs us

Cost per resolution, and what to charge
Gharama kwa kila soko

Each product line's AI spend divided by the markets that line actually **settled**. ⛔ Settled, not merely resolved: a resolved market still inside its objection window has an intact pool and has not been delivered, so counting it would understate the cost.

PRODUCT LINE	AI SPEND	CALLS	SETTLED	COST EACH	CYCLES EACH	SUGGESTED (100%)	UNATTRIBUTED
Polls (generation + Sentinel)	$239.75	320	81	$2.96	0.0296	$5.92 —	$48.00
Up & Down oracle	$2.80	7	0	—	—	— —	$0.00
Help chatbot	$0.1200	6	0	—	—	— —	$0.00
Other / unclassified	$0.00	0	0	—	—	— —	$0.00

⛔ **What 50pick actually charges is a pool commission, not a flat fee.** The "suggested" column is what this AI cost plus a 100% margin would come to per market — it is a cost floor to compare against real commission earned, not a price we charge. TZS 1,000 is the *minimum stake*, not a price.

No USD→TZS rate is set, so every shilling figure reads "—". Set the rate and its date in **Cycle settings** below. A converted figure with no visible rate is a claim nobody can check.

⚠️ **Up & Down AI spend is not the whole cost of an Up & Down round.** One oracle call serves an *observation*, which is shared by every round on that boundary — measured at 2.353 rounds per observation. Rounds settled by the price feed instead of the AI cost nothing in this table at all.

Model mix over this span: claude-sonnet-4-6 100% · claude-haiku-4-5-20251001 0%. A switch between model tiers moves cost per resolution several-fold on its own.

Cost per resolution. Each product line's AI spend divided by the markets that line actually settled.

COLUMN	WHAT IT MEANS
Settled	Markets that finished **and paid out**. A market that has a verdict but is still inside its objection window is not counted — it has not been delivered yet.
Cost each	The AI cost of one settled market on that line.
Suggested	That cost plus the margin setting. **This is a cost floor for comparison, not a price 50pick charges.**
Unattributed	Spend on that line that could not be tied to one specific market — mostly writing polls that were never published. Shown rather than hidden, because leaving it out would make the product look cheaper than it is.

⛔ **50pick charges a pool commission, not a flat fee.**

TZS 1,000 is the **minimum stake** a player can place. It is not a price we charge per market. The "Suggested" column is there to be compared against the commission a market actually earns — never to be quoted to anyone as a price.

Two things this table is not.

The Up & Down line is history. Up & Down stopped using AI to read prices on 3 August 2026 — it uses a price feed now. Its figure describes a period that has ended.

Shillings need a rate. Every TZS figure shows "—" until someone enters the USD→TZS rate and the date it was taken. There is deliberately no default: a converted figure with no rate behind it is a number nobody can check.

Settings

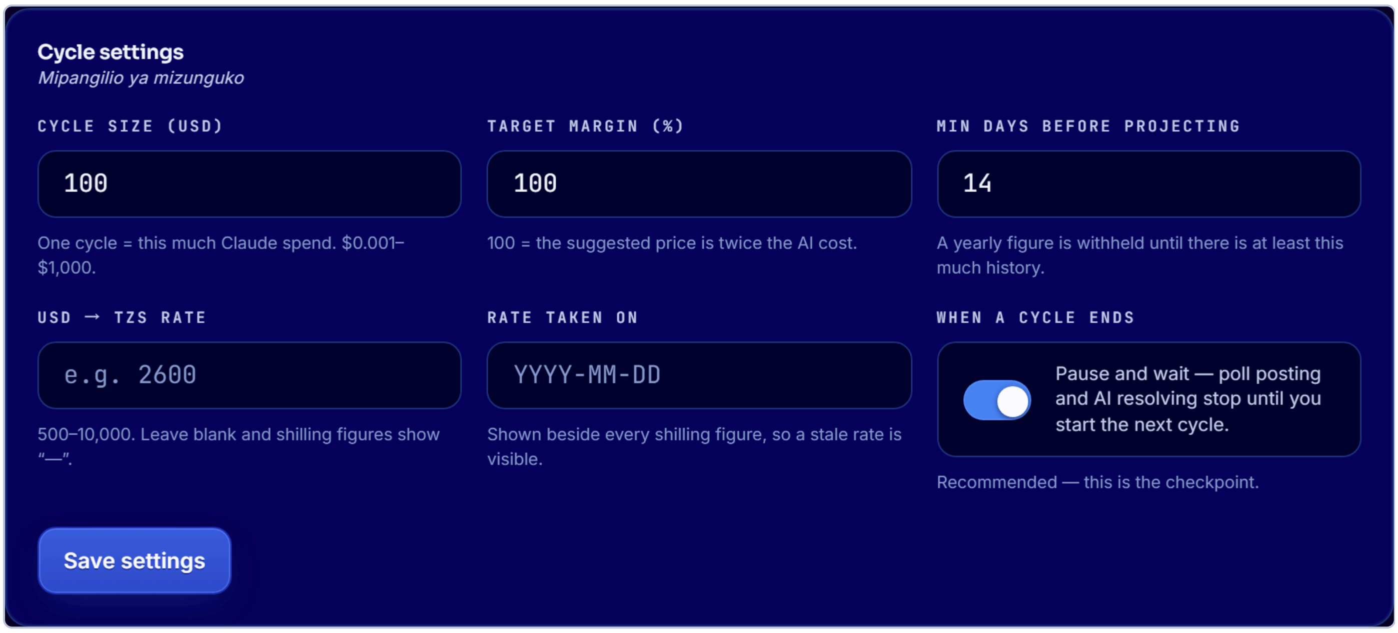

Admin → AI usage → Cycle settings. Changes are confirmed before they save, and every change records who made it and what the value was before.

SETTING	WHAT IT DOES
Cycle size	The amount one cycle represents. **Not retroactive** — cycles already closed keep the size they were opened with, so past counts never move. The new size applies to the next cycle opened.
Target margin	Added to the AI cost to produce the "Suggested" column. 100% means the suggested figure is twice the cost.
Min days before projecting	How much history is needed before a yearly figure is shown at all. Below this, the platform says so instead of guessing.
USD → TZS rate + date	Needed before any shilling figure appears. The rate and its date are printed next to every converted figure, and a rate over 30 days old is flagged.
When a cycle ends	Leave this on **Pause and wait**. Turning it off means the AI never stops — the checkpoint disappears.

Cycles vs. top-up windows — two different things

These are easy to confuse and they are not the same. Both appear on the AI usage page.

	SPEND CYCLE	TOP-UP WINDOW
What it is	An **amount** — $100 of spend	A **period** — since you last bought Anthropic credit
Where	The **Spend cycles** panel	The **Credit budget** panel
What ends it	Spending $100	You, after topping up
Numbering	Never resets — cycle 7 stays cycle 7	Has no number; it just restarts
The button	**Start cycle N**	**New top-up window**

Put simply: **the top-up window is how much credit you have bought. The cycle is the $100 unit that credit is counted in.** Starting a new top-up window does not touch the cycles, and starting a cycle does not add credit.

The routine

1. **Top up Anthropic when the credit is low.** The Credit budget panel emails all admins at 80% and again at 100% of your limit.
   Then press **New top-up window** so the “spent this window” counter starts again.

2. **Set the limit to match what you actually bought.** Otherwise **Funded cycles** is wrong, and it is the figure that tells you how long the money will last.
   Bought $1,000? Set the limit to 1000. That reads as 10 funded cycles.

3. **When the red bar appears, check the two things under Before you press Start, then start the next cycle.**
   Poll posting and AI resolving are stopped until you do. Nothing is lost while it waits.

4. **Once a month, read “Every cycle”.** How long each cycle lasted is the burn rate. If cycles are getting shorter, spending is accelerating.
   Compare it against “Cost per resolution” to see whether markets are costing more, or whether there are simply more of them.

5. **Keep the exchange rate current.** A rate more than 30 days old is flagged on the page.
   A stale rate is a pricing error, not a cosmetic one.

Never

Never turn off “Pause and wait” without a decision from Ali — it removes the checkpoint entirely.

Never quote the “Suggested” column as a price. It is a cost floor for internal comparison.

Never read a figure that says “—” as zero. It means the platform does not have what it needs to answer honestly — usually a missing exchange rate or too little history.